FireFlow

FOR FINANCIAL INSTITUTIONS

by Persistent Labs

From first meeting to live AI.

YOUR DEPLOYMENT JOURNEY

START **Initial meeting**

Agree priority use cases, systems, data boundaries and owners.

Outcome: scope & success criteria

↓

STAGE 1 **Establish staging**

Deploy staging, enable teams and scope the first three agents.

Gate: staging & scope accepted

↓

STAGE 2 **Build & test**

Connect sandbox systems; build agents, permissions and budgets.

Gate: staging tests passed

↓

STAGE 3 **Prepare production**

Set isolation, action controls, audit logs and the infosec plan.

Gate: security & compliance sign-off

↓

STAGE 4 **Limited rollout**

Launch to approved users; monitor and calibrate controls.

Gate: stability & audit review

↓

STAGE 5 **Operate & improve**

Run approved workflows with maintenance, support and usage controls.

Expand through approved projects

Each gate must be accepted before advancing.

WHAT YOU PAY FOR

1. Setup & integration

US$100,000 / Stage 1

One-time foundation setup · invoiced at signature

Stages 2 to 4: integration and rollout quoted separately, with fixed or milestone fees.

2. Maintenance & support

US$20,000 / month

Six-month minimum: **US$120,000 prepaid at Stage 2**; then rolling monthly.

Updates, fixes, supported connectors and agreed support. New development is separate.

3. Compute surcharge

Workload	Frontier APIs	Private models
UIO Customer agents & chats	Not used	25% private only
BPO Internal processes	4% approved tasks	25% approved tasks

Private: institution-hosted, Persistent Labs or approved third party; on-premises or private API.

Frontier: OpenAI / Anthropic APIs. BPO selects by task complexity and data sensitivity; UIO has no frontier fallback.

No defined token cost?

For self-hosted or non-token-priced third-party models, use the **model's market rate per million tokens**, then apply the surcharge.

`Tokens ÷ 1M × market rate × surcharge %`

Agree input/output benchmarks and review cycle. Reference basis only; no repeat charge for institution-paid inference or hardware.

Otherwise: actual FireFlow inference cost after agreed discounts. Monthly billing in arrears; separate cost and surcharge. Infrastructure excluded.

4. Add capabilities through optional projects

Separate statement of work

Governed back-office MCP enablement · new connectors · bespoke UIO/BPO agents · evaluation, migration & training

Scope & price — Fixed fee or capped time & materials → **Build & evaluate** — Implement and test agreed scope → **Approve & release** — Institution acceptance before rollout

↑ Approved additions return to ongoing operation

Institution: infrastructure, APIs, approvals, operations and compliance. FireFlow runs on-premises or in its private cloud.

Persistent Labs: deployment, integration, maintenance and contracted support. Private inference hosting and capacity are separately agreed.

Draft commercial reference. Prices, rates, scope and service targets require agreement. US dollars; taxes excluded. UIO = User Intent Orchestration. BPO = Business Process Orchestration.